

Artificial Intelligence

Local search

Instructors: Fatemeh Mansoori
University of Isfahan



Artificial Intelligence: A Modern Approach*, 3rd Edition, Chapter 4
Most slides have been adapted from CS188, UC Berkeley

Outline

- Local search & optimization algorithms
 - Hill-climbing search
 - Simulated annealing search
 - Local beam search
 - Genetic algorithms



Review: Search

- Find a path with minimum cost from start state to goal state in state space graph
 - Systematic exploration of search space



Local search algorithms

- In many optimization problems, **path** is irrelevant; the goal state **is** the solution
 - state space = set of "complete" configurations;
 - For e.g. in solving n-queen by systematic search, each state is not a complete configuration
- find **configuration satisfying constraints**, e.g., n-queens problem;
- find **optimal configuration**, e.g., travelling salesperson problem



Local search algorithms

- use **iterative improvement** algorithms: keep a single “current” state, try to improve it
- Constant space, suitable for online as well as offline search
- More or less unavoidable if the “state” is yourself (i.e., learning)

Sample problems for local & systematic search

- Path to goal is important
 - Theorem proving
 - Route finding
 - 8-Puzzle
 - Chess
- Goal state itself is important
 - 8 Queens
 - TSP
 - VLSI Layout
 - Job-Shop Scheduling
 - Automatic program generation

Local Search

- Tree search keeps unexplored alternatives on the frontier (ensures completeness)
- Local search: improve a single option until you can’t make it better (no fringe!)
 - New successor function: local changes
- Generally much faster and more memory efficient (but incomplete and suboptimal)

Hill Climbing

- Simple, general idea:
 - Start wherever
 - Repeat: move to the best neighboring state
 - If no neighbors better than current, quit
- What’s about this approach?
 - Complete?
 - Optimal?
 - Benefit ?




n -queens problem

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- Systematic search :
 - What is state-space?
 - What is the size of state space?



Start state



Goal state



Hill-climbing algorithm

- Node contains the state and the value of objective function in that state (not path)
- Search strategy: steepest ascent among immediate neighbors until reaching a peak

function HILL-CLIMBING(problem) **returns** a state

current \leftarrow make-node(problem.initial-state)

loop do

neighbor \leftarrow a highest-valued successor of current


if neighbor.value \leq current.value **then**

return current.state

current \leftarrow neighbor

- Current node is replaced by the best successor (if it is better than current node)
- "Like climbing Everest in thick fog with amnesia"*

Run hill climbing algorithm




n -queens problem

- Modeling the state as a whole configuration
 - States: n queens on board, one per column
 - Successors: move a queen in its column
 - What is the branching factor ?
 - What is objective function?



Heuristic for n -queens problem


- Heuristic value function: number of conflicts



Local search: 8-queens problem

- States: 8 queens on the board, one per column ($8^8 \approx 17$ million)
- Successors(s): all states resulted from s by moving a single queen to another square of the same column ($8 \times 7 = 56$)
- Cost function $h(s)$: number of queen pairs that are attacking each other, directly or indirectly
- Global minimum : $h(s) = 0$

$h(s) = 17$



successors objective values


10	12	14	13	13	14	14	14
14	16	13	15	12	14	12	16
14	16	10	13	15	12	14	14
15	14	14	13	15	13	16	13
14	17	15	13	15	14	16	16
17	16	18	15	15	14	15	16
18	14	15	15	15	14	16	16
14	14	13	17	16	14	16	18

Red: best successors


Hill-climbing search: 8-queens

- Result of hill-climbing in this case...

- A local minimum with $h = 1$




Hill Climbing Quiz




- Starting from X, where do you end up ?
- Starting from Y, where do you end up ?
- Starting from Z, where do you end up ?

Hill-climbing search is greedy

- Greedy local search: considering only one step ahead and select the best successor state (steepest ascent)
- Rapid progress toward a solution
 - Usually quite easy to improve a bad solution



Global and local maxima



Random restarts

- find global optimum

Random sideways moves

- Escape from shoulders
- Loop forever on flat local maxima

Random-restart hill climbing


- Hill climbing is incomplete
 - Getting stuck on local max
- Hill climbing with random restart
 - **while** state \neq goal **do**
 - run hill-climbing search from a random initial state
- p : probability of success in each hill-climbing search
 - Expected no of restarts = $1/p$
- Reasonable solution can be usually obtained after a small no of restarts

hill climbing variants

- Hill climbing : generate all successors and choose the best one
- Stochastic hill climbing : Randomly chooses among the available uphill moves according to the steepness of these moves
 - $P(S')$ is an increasing function of $h(s') - h(s)$
- First-choice hill climbing: generating successors randomly until one better than the current state is found
 - Good when number of successors is high


Sideways move

- plateau may be a shoulder so keep going sideways moves when there is no uphill move
 - Problem: infinite loop where flat local max
 - Solution: upper bound on the number of consecutive sideways moves



Hill-climbing on the 8-queens problem

- No sideways moves:
 - Succeeds w/ prob. 0.14
 - Expected number of restart to find the answer ? $1/p = 7$
 - Average number of moves per trial (cost of finding solution):
 - 4 when succeeding, 3 when getting stuck
 - Expected total number of moves needed: $3(1-p)/p + 4 \approx 22$ moves
- Allowing 100 sideways moves:
 - Succeeds w/ prob. 0.94
 - Average number of moves per trial:
 - 21 when succeeding, 65 when getting stuck
 - Expected total number of moves needed: $65(1-p)/p + 21 \approx 25$ moves



Simulated annealing

- Resembles the annealing process used to cool metals slowly to reach an ordered (low-energy) state
- Basic idea:
 - Allow "bad" moves occasionally, depending on "temperature"
 - High temperature \Rightarrow more bad moves allowed, shake the system out of its local minimum
 - Gradually reduce temperature according to some schedule
 - Sounds pretty flaky, doesn't it?

Simulated annealing algorithm

```


function SIMULATED-ANNEALING(problem,schedule) returns a state
  inputs : problem : a problem
          schedule : a mapping from time to "temperature"
  current  $\leftarrow$  problem.initial-state
  for t = 1 to  $\infty$  do
    T  $\leftarrow$  schedule(t) (T is a temperature, controlling probability of
      downward steps)
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  next.value - current.value
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
    
```

- Theoretical guarantee:
 - If T decreased slowly enough, will converge to optimal state
 - Convergence can be guaranteed if at each step, T drops no more quickly than $C/\log n$, $C = \text{constant}$, $n = \#$ of steps so far



Simulated Annealing

- Is this convergence an interesting guarantee?
- Sounds like magic, but reality is reality:
 - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all in a row
 - "Slowly enough" may mean exponentially slowly
 - Random restart hill climbing also converges to optimal state...
- Simulated annealing and its relatives are a key workhorse in VLSI layout and other optimal configuration problems




Local beam search

- Basic idea:
 - K copies of a local search algorithm, initialized randomly
- For each iteration
 - Generate ALL successors from K current states
 - If any of successors is goal \rightarrow finished
 - Choose best K of these to be the new current states

Or, K chosen randomly with a bias towards good ones



Local beam search example ($k=4$)



Local beam search


- Why is this different from K local searches in parallel?
 - The searches **communicate!** "Come over here, the grass is greener!"
- Problem: quite often, all k states end up on same local hill
 - Idea: Stochastic beam search
 - Choose k successors randomly, biased towards good ones
- What other well-known algorithm does this remind you of?
 - Evolution!

الف) در صورت استفاده از first choice hill climbing به کدام همسایه می رویم. (ترتیب ساخته شدن همسایه به ترتیب شماره هاست)


ب) در صورت استفاده از stochastic hill climbing کدام حالت ها ممکن است انتخاب شوند، و احتمال انتخاب هر کدام چقدر است؟

ج) در صورت استفاده از hill climbing معمولی کدام حالت انتخاب می شود؟ و آیا از این حالت می توانیم به هدف برسیم یا یک ماکزیمم محلی است؟

د) در صورت استفاده از الگوریتم beam search اگر $k=3$ باشد، کدام حالت یا حالت ها انتخاب می شوند؟




Genetic algorithms



- Genetic algorithms use a natural selection metaphor
 - Resample K individuals at each step (selection) weighted by fitness function
 - Combine by pairwise crossover operators, plus mutation to give variety


Example: N-Queens



- Does crossover make sense here?
- What would mutation be?
- What would a good fitness function be?


Chromosome & fitness: 8-queens

- Describe the individual (or state) as a string



- Fitness function: number of non-attacking pairs of queens
 - 24 for above figure


Genetic operators: Selection



- Fitness function: number of non-attacking pairs of queens
 - min = 0, max = $8 \times 7/2 = 28$
- Reproduction rate(i) = $\text{fitness}(i) / \sum_{k=1}^n \text{fitness}(k)$
 - e.g., $24 / (24 + 23 + 20 + 11) = 31\%$


Genetic operators: Cross-over

- To select some part of the state from one parent and the rest from another.

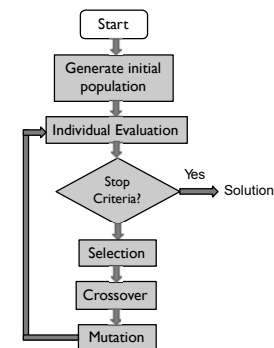


Genetic operators: Mutation

- change a small part of one state with a small probability



A Genetic algorithm diagram



Genetic algorithm properties

- Genetic algorithm is a variant of “stochastic beam search”
- Why does a genetic algorithm usually take large steps in earlier generations and smaller steps later?
 - Initially, population individuals are diverse
 - Cross-over operation on different parent states can produce a state long a way from both parents
 - More similar individuals gradually appear in the population
- Useful on some set of problems but no convincing evidence that GAs are better than hill-climbing w/random restarts in general

Local search vs. systematic search

	Systematic search	Local search
Solution	Path from initial state to the goal	Solution state itself
Method	Systematically trying different paths from an initial state	Keeping a single or more “current” states and trying to improve them
State space	Usually incremental	Complete configuration
Memory	Usually very high	Usually very little (constant)
Time	Finding optimal solutions in small state spaces	Finding reasonable solutions in large or infinite (continuous) state spaces
Scope	Search	Search & optimization problems

38

Question

- Give the name of the algorithm that results from each of the following special cases
 - Local beam search with $k = 1$
 - Local beam search with one initial state and no limit on the number of states retained.
 - Simulated annealing with $T = 0$ at all times (and omitting the termination test).
 - Simulated annealing with $T = \infty$ at all times.
 - Genetic algorithm with population size $N = 1$.